

# DXRAM

## Log + Recovery

Florian Klein

Institut für Informatik  
Abteilung Betriebssysteme  
Heinrich-Heine-Universität Düsseldorf

11.12.2012



# Outline

- 1 Data-Structure
- 2 Log
- 3 Recovery
- 4 Conclusion

# Data-Structure

## ChunkInformation

<i>Field</i>	<i>Type</i>	<i>Description</i>
m_chunkID	ID	unique Chunk-Identifier
m_version	long	version of the Chunk
m_flags	int	flags (present, foreign)
m_nodeID	NodeID	unique identifier of the hosting Node

## Chunk extends ChunkInformation

<i>Field</i>	<i>Type</i>	<i>Description</i>
m_data	ByteBuffer	binary data

# Outline

## 1 Data-Structure

## 2 Log

- Current Log-Layout
- Update process

## 3 Recovery

## 4 Conclusion

# Types

## Block

- stores a Chunk
- default 1 MB

## Segment

- merges a number of Chunks
- backup entity
- default 64 MB

# Meta-Data I

## Segment-Header

- Segment-Index (4 Byte)
- Flags (4 Byte)
  - ▶ empty
  - ▶ full
  - ▶ foreign
  - ▶ number of free Blocks
- Bitmap (1 Bit per Block in the Segment)

## Bitmap (64 MB Segments and 1 MB Blocks)

⇒ 64 Blocks per Segment ⇒ 8 Byte Bitmap

## Bitmap (64 MB Segments and 1 KB Blocks)

⇒ 65536 Blocks per Segment ⇒ 8192 Byte Bitmap

# Meta-Data II

## Log-Header

### Hash-Table with one entry per Block

Bytes	Content	Type	
1	valid-flag	boolean	
4	ID length	ID	ChunkInformation
20	ID bytes		
8	version	long	
4	flags	int	
4	host length	NodeID	
15	host		
2	port		
8	address	long	StorageInformation
4	size	int	
70			

# Layout

## Initial situation

NodeID: 10

Segmentsize: 64 MB

Blocksize: 1 MB

0x00000000

Hash-Table						
Valid	ID	Version	Flags	NodeID	Address	Size

0x006F4C30

Segment 1							
Bitmap: 0000...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: empty							

0x046F4C50

Segment 2							
Bitmap: 0000...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: empty							

Chunk 13	
Version: 1	Data 2500 KB
Flags: present	
NodeID: 10	

Chunk 24	
Version: 1	Data 62000 KB
Flags: present	
NodeID: 10	

Chunk 17	
Version: 1	Data 2000 KB
Flags: foreign	
NodeID: 21	



# Layout

## Inserting Chunk 13

NodeID: 10

Segmentsize: 64 MB

Blocksize: 1 MB

0x00000000

Hash-Table							
Valid	ID	Version	Flags	NodeID	Address	Size	
1	13	1	present	10	0x006F4C40	2500 KB	

0x006F4C30

Segment 1							
Bitmap: 1110...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags:							

0x046F4C50

Segment 2							
Bitmap: 0000...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: empty							

Chunk 13	
Version: 1	Data 2500 KB
Flags: present	
NodeID: 10	

Chunk 24	
Version: 1	Data 62000 KB
Flags: present	
NodeID: 10	

Chunk 17	
Version: 1	Data 2000 KB
Flags: foreign	
NodeID: 21	

# Layout

## Inserting Chunk 24

NodeID: 10

Segmentsize: 64 MB

Blocksize: 1 MB

0x00000000

Hash-Table							
Valid	ID	Version	Flags	NodeID	Address	Size	
1	13	1	present	10	0x006F4C40	2500 KB	
1	24	3	present	10	0x006F5840	62000 KB	

0x006F4C30

Segment 1							
Bitmap: 1111...11	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: full							

0x046F4C50

Segment 2							
Bitmap: 0000...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: empty							

Chunk 13	
Version: 1	Data 2500 KB
Flags: present	
NodeID: 10	

Chunk 24	
Version: 3	Data 62000 KB
Flags: present	
NodeID: 10	

Chunk 17	
Version: 1	Data 2000 KB
Flags: foreign	
NodeID: 21	

# Layout

## Inserting Chunk 17

NodeID: 10

Segmentsize: 64 MB

Blocksize: 1 MB

0x00000000

Hash-Table							
Valid	ID	Version	Flags	NodeID	Address	Size	
1	13	1	present	10	0x006F4C40	2500 KB	
1	24	3	present	10	0x006F5840	62000 KB	
1	17	1	foreign	21	0x046F4C60	2000 KB	

0x006F4C30

Segment 1							
Bitmap: 1111...11	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: full							

0x046F4C50

Segment 2							
Bitmap: 1100...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: foreign							

Chunk 13	
Version: 1	Data 2500 KB
Flags: present	
NodeID: 10	

Chunk 24	
Version: 3	Data 62000 KB
Flags: present	
NodeID: 10	

Chunk 17	
Version: 1	Data 2000 KB
Flags: foreign	
NodeID: 21	

## Final situation

NodeID: 10

Segmentsize: 64 MB

Blocksize: 1 MB

0x00000000

Hash-Table						
Valid	ID	Version	Flags	NodeID	Address	Size
1	13	1	present	10	0x006F4C40	2500 KB
1	24	3	present	10	0x006F5840	62000 KB
1	17	1	foreign	21	0x046F4C60	2000 KB

0x006F4C30

Segment 1							
Bitmap: 1111...11	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: full							

0x046F4C50

Segment 2							
Bitmap: 1100...00	Block 1	Block 2	Block 3	Block 4	...	Block 63	Block 64
Flags: foreign							

Chunk 13	
Version: 1	Data 2500 KB
Flags: present	
NodeID: 10	

Chunk 24	
Version: 3	Data 62000 KB
Flags: present	
NodeID: 10	

Chunk 17	
Version: 1	Data 2000 KB
Flags: foreign	
NodeID: 21	

# Overhead

## Log-Header

- 70 Bytes per Block

## Segment-Header

- 1 Bit per Block
- 8 Bytes per Segment

# Overhead examples

## 100 GB Log-File, 64 MB Segments and 1 MB Blocks

⇒ 1.600 Segements each with 64 Blocks

⇒ 102.400 Blocks

$$\begin{aligned} &1.600 * 8B + 102.400 * 1Bit + 102.400 * 70B \\ &= 12.800B + 12.800B + 7.168.000B = 7.193.600B \\ &\approx 7MB \end{aligned}$$

## 100 GB Log-File, 64 MB Segments and 1 KB Blocks


⇒ 1.600 Segements each with 65536 Blocks

⇒ 104.857.600 Blocks

$$\begin{aligned} &1.600 * 8B + 104.857.600 * 1Bit + 104.857.600 * 70B \\ &= 12.800B + 13.107.200B + 7.340.032.000B = 7.353.152.000B \\ &\approx 7GB \end{aligned}$$


# Update process (Overview)

## Initial situation



# Update process (Overview)


## 1. Step






# Update process (Overview)

## 2. Step




# Update process (Overview)

## 3. Step




# Update process (Overview)

## 4. Step




# Update process (Overview)

## 5. Step




# Update process (Overview)

## 6. Step




# Update process (Overview)

## 7. Step




# Update process (Overview)

## 8. Step



# Update process (Overview)

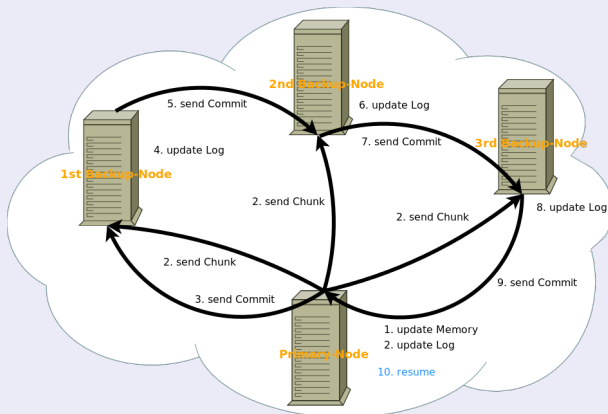
## 9. Step





# Update process (Overview)

## 10. Step



# Update process I

## Memory (Primary-Node)

- (a) set dirty-bit
- (b) update chunk in-place
- (c) increment version
- (d) clear dirty-bit

## Log (Primary-Node)

- (a) send chunk to the backup nodes
- (b) update chunk
  - Solution 1: in-place
  - Solution 2: copy-on-write (RAMCloud)
  - Solution 3: twins
- (c) send Commit-Message to 1<sup>st</sup> Backup-Node
- (d) wait for Commit-Message from 3<sup>rd</sup> Backup-Node

# Update process II

## Log (Backup-Node)

- (a) receive chunk from Primary-Node
- (b) buffer chunk
- (c) wait for Commit-Message from predecessor
- (d) update chunk (like Primary-Node)
- (e) send Commit-Message to successor

## Problems

- Primary-Node crashes before sending the Commit-Message
- Backup-Node crashes

# Outline

1 Data-Structure

2 Log


3 Recovery

- Peers and Super-Peers
- Recovery process

4 Conclusion

# Super-Peer Overlay

## Overview



# Tasks

## Peer


- stores Chunks in RAM
- logs own Chunks on SSD
- backups Chunks of other nodes on SSD

## Super-Peer

- also a Peer
- monitors its Peers (Heart-Beat)
- handles Meta-Data of Chunks (DHT-Node)

# Recovery process (Overview)


## Initial situation



- Super-Peer is connected with its Peers

# Recovery process (Overview)

## Node failure




- Node 27 crashes



# Recovery process (Overview)


## Broadcast



- Super-Peer detects and broadcasts node failure
- Super-Peer becomes the Recovery Coordinator

# Recovery process (Overview)

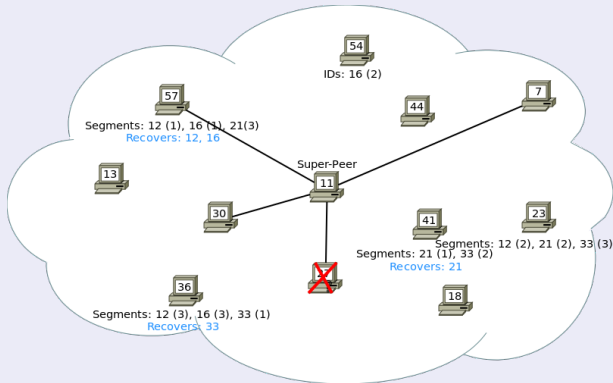
## Info-Message



- every node deletes cached Meta-Data of the failed node
- every Backup-Node of the failed node sends information about the backedup data to the Recovery Coordinator
- Recovery Coordinator gathers this information

# Recovery process (Overview)

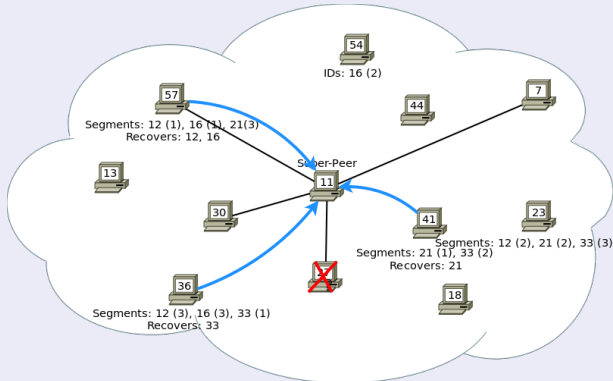
## Chunk recovery



- the 1<sup>st</sup> Backup-Node of each segment recovers the segment

# Recovery process (Overview)


## Commit-Message



- the 1<sup>st</sup> Backup-Node of each segment send the result to the Recovery Coordinator

# Recovery process (Overview)

## Final situation



- Recovery Coordinator updates Meta-Data for the recovered Segments/Chunks
- the failed node is removed
- an application callback can be executed (optional)

# Conclusion

## Log

- little overhead for objects  $\geq 1MB$  (e.g. Map & Reduce)
  - huge overhead for Objects  $\leq 1KB$  (e.g. Facebook)
- ⇒ a new log layout have to be designed for small objects
- the three update solutions must be compared

## Recovery

- distributed recovery dependent on segment distribution
  - ▶ high distribution allows fast recovery
  - ▶ low distribution preserves locality
- ignores workload of the Backup-Nodes
- Problem: Not all data could be locally recovered (IO failure or insufficient memory space)
- Problem: Realisation of the application callback

Questions?